

# 50\_STARTUPS

```
In [74]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.graphics.regressionplots import influence_plot
import statsmodels.formula.api as snf
import numpy as np
import statsmodels.api as sm
```

```
In [53]: st=pd.read_csv("C:\\Users\\nishi\\Downloads\\50_Startups.csv")
st
```

```
Out[53]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92

	R&D Spend	Administration	Marketing Spend	State	Profit
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

In [54]: `st.columns = st.columns.str.replace(' ', '_')`

In [55]: `st.head()`

Out[55]:

	R&D_Spend	Administration	Marketing_Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

In [56]: `st['R&D_Spend'].mean()`

Out[56]: 73721.61559999999

In [57]: `st['Administration'].mean()`

Out[57]: 121344.63959999995

```
In [58]: st['Marketing_Spend'].mean()
```

```
Out[58]: 211025.097800000005
```

```
In [59]: st['Profit'].mean()
```

```
Out[59]: 112012.639200000002
```

```
In [60]: st.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 50 entries, 0 to 49  
Data columns (total 5 columns):  
#   Column              Non-Null Count  Dtype  
---  ---  
0   R&D_Spend           50 non-null     float64  
1   Administration      50 non-null     float64  
2   Marketing_Spend     50 non-null     float64  
3   State               50 non-null     object  
4   Profit              50 non-null     float64  
dtypes: float64(4), object(1)  
memory usage: 2.1+ KB
```

```
In [61]: st1=st.rename({'R&D_Spend':'rdSpend','Administration':'Adm','Marketing_Spend':'Ma
st1
```

```
Out[61]:
```

	rdSpend	Adm	Mark	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92

	rdSpend	Adm	Mark	State	Profit
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

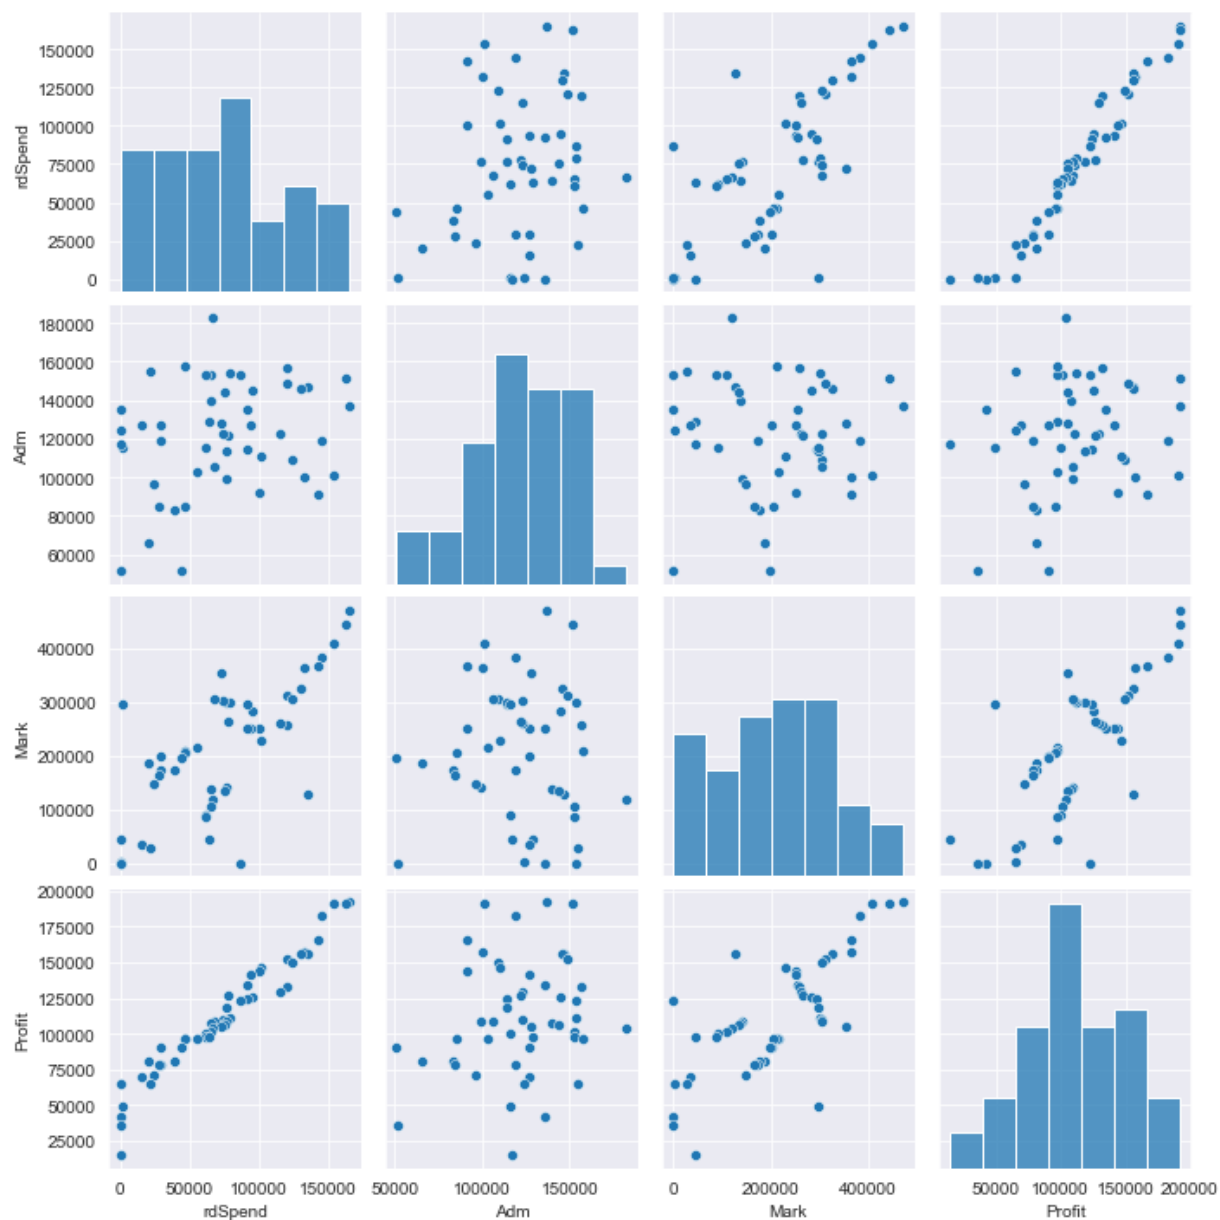
In [62]: *# finding the correlaation*  
st1.corr()

Out[62]:

	rdSpend	Adm	Mark	Profit
rdSpend	1.000000	0.241955	0.724248	0.972900
Adm	0.241955	1.000000	-0.032154	0.200717
Mark	0.724248	-0.032154	1.000000	0.747766
Profit	0.972900	0.200717	0.747766	1.000000

```
In [63]: sns.set_style(style='darkgrid')
sns.pairplot(st1)
```

```
Out[63]: <seaborn.axisgrid.PairGrid at 0x1d0699d99a0>
```



```
In [64]: # check for null values
st1.isna().sum()
```

```
Out[64]: rdSpend      0
         Adm          0
         Mark         0
         State        0
         Profit       0
         dtype: int64
```

```
In [65]: #Build the model
import statsmodels.formula.api as snf
model=snf.ols('Profit~Adm+Mark+rdSpend',data=st1).fit()
```



```
In [13]: model.fittedvalues
```

```
Out[13]: 0      192521.252890
          1      189156.768232
          2      182147.279096
          3      173696.700026
          4      172139.514183
          5      163580.780571
          6      158114.096669
          7      160021.363048
          8      151741.699699
          9      154884.684110
         10      135509.016367
         11      135573.712961
         12      129138.054182
         13      127487.991663
         14      149548.646335
         15      146235.159985
         16      116915.405401
         17      130192.447208
         18      129014.226806
         19      115635.216367
         20      116639.669231
         21      117319.451640
         22      114706.981717
         23      109996.615221
         24      113362.966113
         25      102237.725065
         26      110600.575350
         27      114408.071457
         28      101660.026005
         29      101794.983452
         30       99452.372936
         31       97687.856276
         32       99001.328985
         33       97915.007805
         34       89039.273741
         35       90511.599568
         36       75286.174585
         37       89619.537708
         38       69697.430648
         39       83729.011977
         40       74815.953991
         41       74802.556239
         42       70620.411821
         43       60167.039963
         44       64611.354916
         45       47650.649687
         46       56166.206853
         47       46490.588983
         48       49171.388158
         49       48215.134111
dtype: float64
```

```
In [66]: # beta coefficients
model.params
```

```
Out[66]: Intercept    50122.192990
Adm                 -0.026816
Mark                 0.027228
rdSpend             0.805715
dtype: float64
```

```
In [67]: # t and p values
(model.pvalues, "\n", model.tvalues)
```

```
Out[67]: (Intercept    1.057379e-09
Adm                 6.017551e-01
Mark                 1.047168e-01
rdSpend             2.634968e-22
dtype: float64,
'\n',
Intercept    7.626218
Adm          -0.525507
Mark          1.655077
rdSpend      17.846374
dtype: float64)
```

```
In [68]: # r squared values
(model.rsquared, model.rsquared_adj)
```

```
Out[68]: (0.9507459940683246, 0.9475337762901719)
```

```
In [ ]: #NOW FOR THE MULTI COLLINEARITY PROBLEM WE NEED TO FIND OUT THE Simple REGRESSION
#IT IS A GOOD PRACTICE
```

```
In [69]: m1_v=snf.ols('Profit~Adm', data=st1).fit()
#p and the t values
print(m1_v.pvalues, "\n", m1_v.tvalues)
```

```
Intercept    0.003824
Adm           0.162217
dtype: float64
Intercept    3.040044
Adm           1.419493
dtype: float64
```

```
In [70]: m1_v=snf.ols('Profit~Mark',data=st1).fit()
#p and the t values
print(m1_v.pvalues,"\n",m1_v.tvalues)
```

```
Intercept    4.294735e-10
Mark          4.381073e-10
dtype: float64
Intercept      7.808356
Mark           7.802657
dtype: float64
```

```
In [71]: m1_v=snf.ols('rdSpend~Adm+Mark',data=st1).fit()
#p and the t values
print(m1_v.pvalues,"\n",m1_v.tvalues)
```

```
Intercept    7.713519e-02
Adm           6.322635e-03
Mark          3.724804e-10
dtype: float64
Intercept    -1.807194
Adm           2.858702
Mark          7.889568
dtype: float64
```

```
In [73]: rsq_rd=snf.ols("rdSpend~Adm+Mark",data=st1).fit().rsquared
vif_rd=1/(1-rsq_rd)

rsq_adm=snf.ols("Adm~rdSpend+Mark",data=st1).fit().rsquared
vif_adm=1/(1-rsq_adm)


rsq_mark=snf.ols("Mark~Adm+rdSpend",data=st1).fit().rsquared
vif_mark=1/(1-rsq_mark)

# Putting the values in Dataframe format
d1={'Variables':['rdSpend','Administration','Marketing_Spend'],'Vif':[vif_rd,vif_
Vif_df=pd.DataFrame(d1)
Vif_df
```

```
Out[73]:
```

	Variables	Vif
0	rdSpend	2.468903
1	Administration	1.175091
2	Marketing_Spend	2.326773

```
In [75]: sm.qqplot(model.resid,line='q')
plt.title("Normal Q-Q plot of residuals")
plt.show()
```




```
In [76]: list(np.where(model.resid<-20000))
```

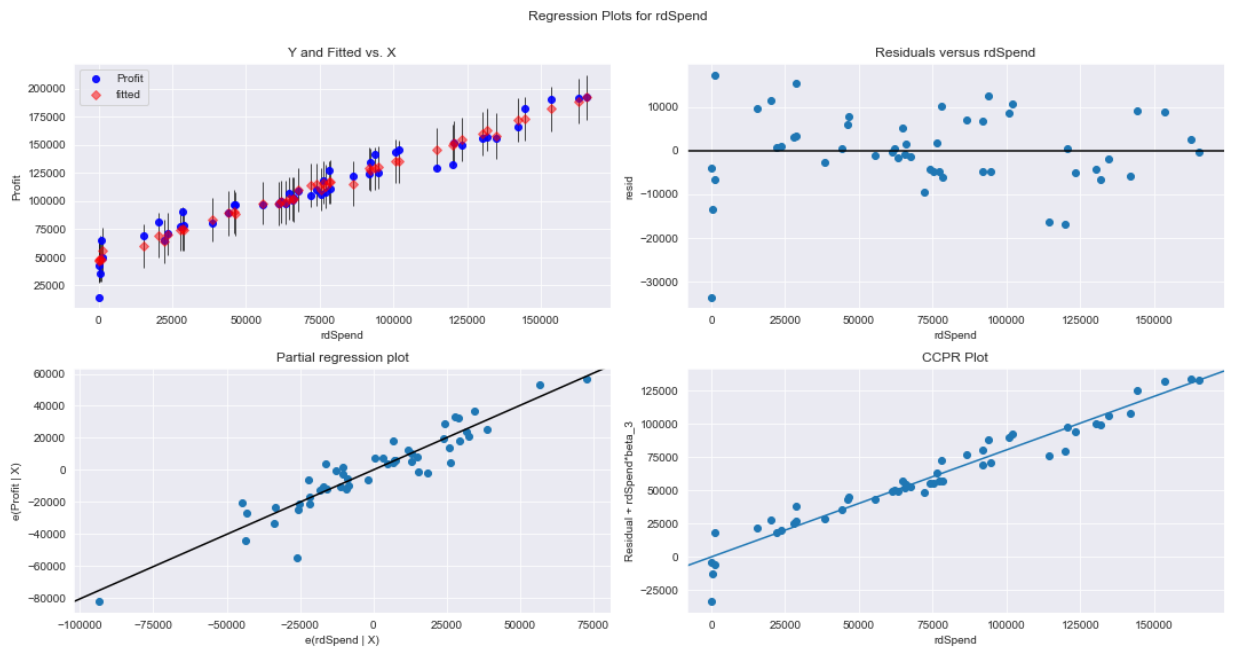
```
Out[76]: [array([49], dtype=int64)]
```

```
In [77]: def standard_values(vals) : return (vals-vals.mean())/vals.std()
```

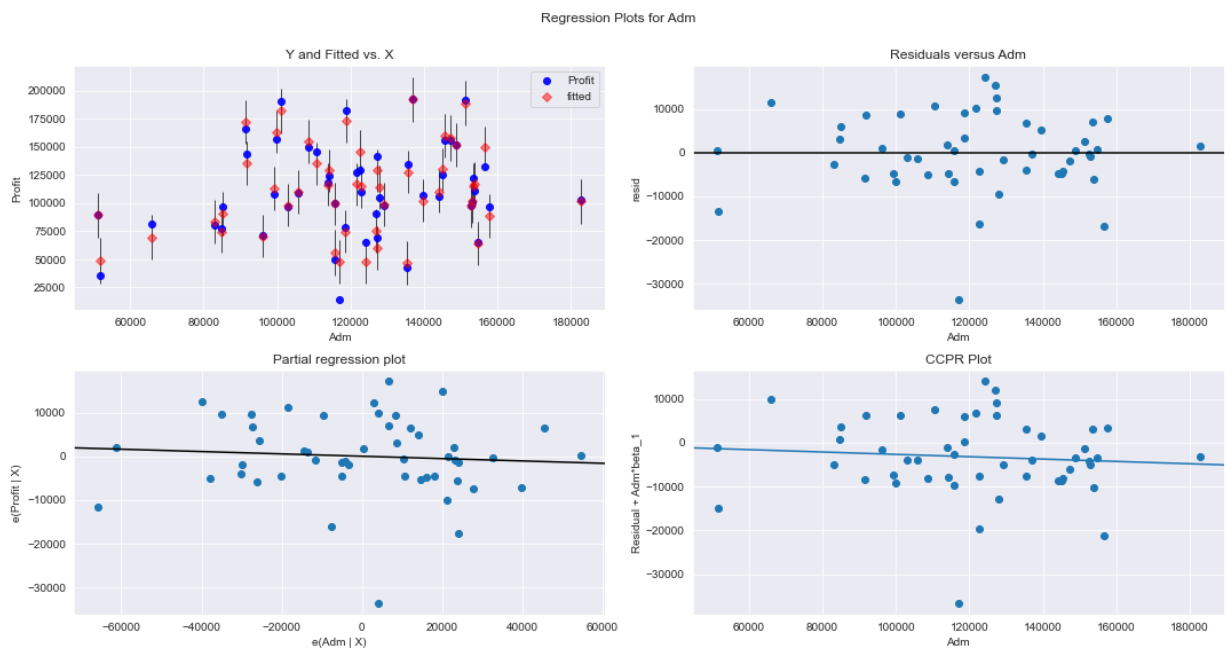
```
In [78]: plt.scatter(standard_values(model.fittedvalues),standard_values(model.resid))
plt.title('Residual Plot')
plt.xlabel('standardized fitted values')
plt.ylabel('standardized residual values')
plt.show()
```



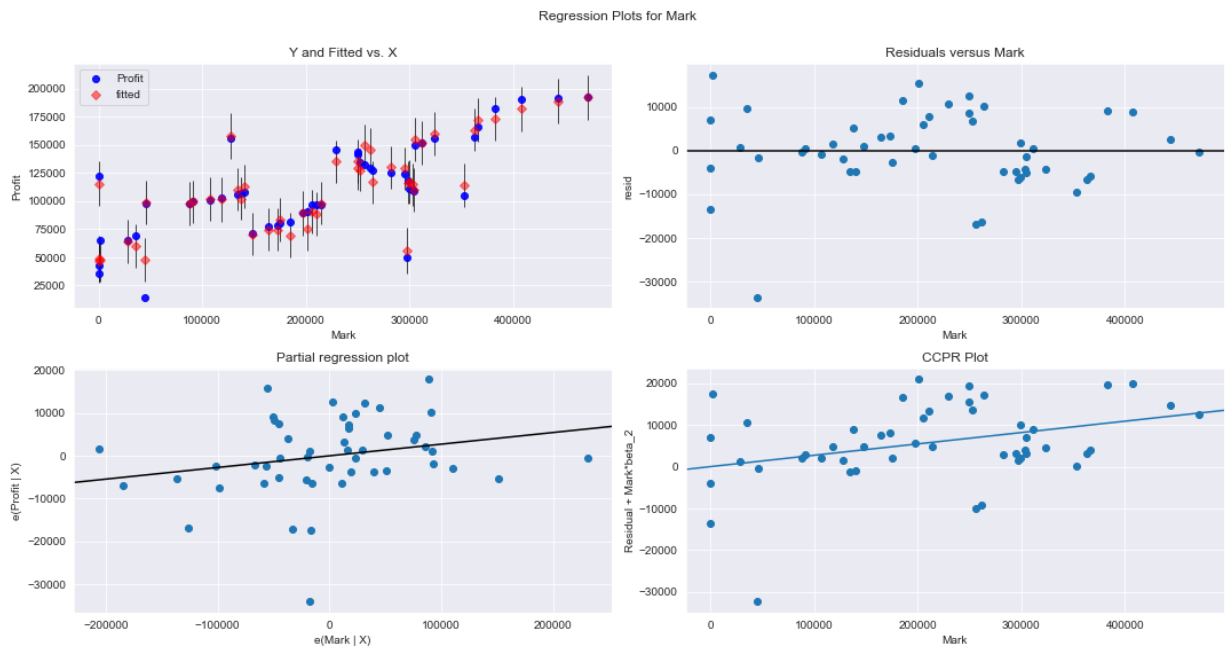
```
In [79]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'rdSpend',fig=fig)
plt.show()
```



```
In [80]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'Adm',fig=fig)
plt.show()
```




```
In [81]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'Mark',fig=fig)
plt.show()
```



```
In [82]: (c,_)=model.get_influence().cooks_distance
c
```

```
Out[82]: array([3.21825244e-05, 3.27591036e-03, 3.23842699e-02, 2.17206555e-02,
 1.44833032e-02, 1.17158463e-02, 2.91766303e-03, 3.56513444e-03,
 4.04303948e-05, 4.86758017e-03, 1.51064757e-02, 1.63564959e-02,
 1.15516625e-02, 4.01422811e-03, 6.12934253e-02, 3.40013448e-02,
 8.33556413e-03, 3.30534399e-03, 2.16819303e-03, 4.07440577e-02,
 4.25137222e-04, 1.09844352e-02, 2.91768000e-03, 2.76030254e-04,
 5.04643588e-03, 3.00074623e-03, 3.41957068e-03, 2.98396413e-02,
 1.31590664e-03, 1.25992620e-04, 4.18505125e-05, 9.27434786e-06,
 7.08656521e-04, 1.28122674e-04, 2.09815032e-02, 6.69508674e-03,
 5.55314705e-02, 6.55050578e-05, 5.61547311e-02, 1.54279607e-03,
 1.84850929e-03, 1.97578066e-03, 1.36089280e-04, 2.05553171e-02,
 1.23156041e-04, 9.03234206e-02, 5.45303387e-02, 5.33885616e-03,
 1.90527441e-01, 2.88082293e-01])
```


```
In [83]: fig=plt.figure(figsize=(20,7))
plt.stem(np.arange(len(st1)),np.round(c,5))
plt.xlabel('Row Index')
plt.ylabel('Cooks Distance')
plt.show()
```



```
In [84]: np.argmax(c) , np.max(c)
```

```
Out[84]: (49, 0.28808229275432634)
```

```
In [85]: influence_plot(model)
plt.show()
```



```
In [87]: k=st1.shape[1]
n=st1.shape[0]
leverage_cutoff = (3*(k+1))/n
leverage_cutoff
```

```
Out[87]: 0.36
```



```
In [88]: st1[st1.index.isin([49])]
```

```
Out[88]:
```

	rdSpend	Adm	Mark	State	Profit
49	0.0	116983.8	45173.06	California	14681.4

```
In [89]: st2=st1.drop(st1.index[[49]],axis=0).reset_index(drop=True)
st2
```

```
Out[89]:
```

	rdSpend	Adm	Mark	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92

	rdSpend	Adm	Mark	State	Profit
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41

```
In [90]: while np.max(c)>0.5 :
          model=snf.ols("Profit~rdSpend+Adm+Mark",data=st2).fit()
          (c,_)=model.get_influence().cooks_distance
          c
          np.argmax(c) , np.max(c)
          st2=st2.drop(st2.index[[np.argmax(c)]],axis=0).reset_index(drop=True)
          data2
        else:
          final_model=snf.ols("Profit~rdSpend+Adm+Mark",data=st2).fit()
          final_model.rsquared , final_model.aic
          print("Thus model accuracy is improved to",final_model.rsquared)
```

Thus model accuracy is improved to 0.9613162435129847

```
In [91]: final_model.rsquared
```

Out[91]: 0.9613162435129847

In [92]:

st2

Out[92]:

	rdSpend	Adm	Mark	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92

	rdSpend	Adm	Mark	State	Profit
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41

In [93]: `new_data=pd.DataFrame({'rdSpend':60000,"Adm":70000,"Mark":130000},index=[0])`  
`new_data`

Out[93]:

	rdSpend	Adm	Mark
0	60000	70000	130000

In [94]: `final_model.predict(new_data)`

Out[94]: 0    101088.827113  
dtype: float64

```
In [95]: pred_y=final_model.predict(st2)
pred_y
```

```
Out[95]: 0      190716.676999
1      187537.122227
2      180575.526396
3      172461.144642
4      170863.486721
5      162582.583177
6      157741.338633
7      159347.735318
8      151328.826941
9      154236.846778
10     135507.792682
11     135472.855621
12     129355.599449
13     127780.129139
14     149295.404796
15     145937.941975
16     117437.627921
17     130408.626295
18     129129.234457
19     116641.003121
20     117097.731866
21     117911.019038
22     115248.217796
23     110603.139045
24     114051.073877
25     103398.054385
26     111547.638935
27     114916.165026
28     103027.229434
29     103057.621761
30     100656.410227
31      99088.213693
32     100325.741335
33      98962.303136
34      90552.307809
35      91709.288672
36      77080.554255
37      90722.503244
38      71433.021956
39      85147.375646
40      76625.510303
41      76492.145175
42      72492.394974
43      62592.049718
44      67025.731107
45      50457.297206
46      58338.443625
47      49375.776655
48      51658.096812
dtype: float64
```

```
In [96]: d2={'Prep_Models':['Model','Final_Model'],'Rsquared':[model.rsquared,final_model.rsquared]}
table=pd.DataFrame(d2)
table
```

```
Out[96]:
```

	Prep_Models	Rsquared
0	Model	0.950746
1	Final_Model	0.961316

## ToyotaCorolla

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.graphics.regressionplots import influence_plot
import statsmodels.formula.api as smf
import numpy as np
import statsmodels.api as sm
```

```
In [2]: cars=pd.read_csv("C:\\Users\\nishi\\Desktop\\Assignments\\Multi_Linear_Regressor\\ToyotaCorolla.csv")
```

In [16]: cars

Out[16]:

	Id	Model	Price	Age_08_04	KM	HP	Doors	Cylinders	Gears	Weight
0	1	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13500	23	46986	90	3	4	5	1165
1	2	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13750	23	72937	90	3	4	5	1165
2	3	ÉTOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13950	24	41711	90	3	4	5	1165
3	4	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	14950	26	48000	90	3	4	5	1165
4	5	TOYOTA Corolla 2.0 D4D HATCHB SOL 2/3-Doors	13750	30	38500	90	3	4	5	1170
...	...	...	...	...	...	...	...	...	...	...
1431	1438	TOYOTA Corolla 1.3 16V HATCHB G6 2/3-Doors	7500	69	20544	86	3	4	5	1025
1432	1439	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	10845	72	19000	86	3	4	5	1015
1433	1440	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	8500	71	17016	86	3	4	5	1015
1434	1441	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	7250	70	16916	86	3	4	5	1015
1435	1442	TOYOTA Corolla 1.6 LB LINEA TERRA 4/5-Doors	6950	76	1	110	5	4	5	1114

1436 rows × 10 columns

In [19]: cars1=cars.drop("Cylinders",axis=1) # We drop the Cylinders column as it has no v



In [20]: cars1

Out[20]:

	Id	Model	Price	Age_08_04	KM	HP	Doors	Gears	Weight
0	1	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13500	23	46986	90	3	5	1165
1	2	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13750	23	72937	90	3	5	1165
2	3	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	13950	24	41711	90	3	5	1165
3	4	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3-Doors	14950	26	48000	90	3	5	1165
4	5	TOYOTA Corolla 2.0 D4D HATCHB SOL 2/3-Doors	13750	30	38500	90	3	5	1170
...	...	...	...	...	...	...	...	...	...
1431	1438	TOYOTA Corolla 1.3 16V HATCHB G6 2/3-Doors	7500	69	20544	86	3	5	1025
1432	1439	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	10845	72	19000	86	3	5	1015
1433	1440	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	8500	71	17016	86	3	5	1015
1434	1441	TOYOTA Corolla 1.3 16V HATCHB LINEA TERRA 2/3-...	7250	70	16916	86	3	5	1015
1435	1442	TOYOTA Corolla 1.6 LB LINEA TERRA 4/5-Doors	6950	76	1	110	5	5	1114

1436 rows × 9 columns

In [21]: cars1["Price"].mean()

Out[21]: 10730.824512534818

In [22]: cars1["HP"].mean()

Out[22]: 101.50208913649026

In [23]: cars1["Doors"].mean()

Out[23]: 4.0334261838440115

In [24]: cars1["Gears"].mean()

Out[24]: 5.0264623955431755

```
In [25]: cars1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   Id              1436 non-null   int64  
 1   Model           1436 non-null   object  
 2   Price           1436 non-null   int64  
 3   Age_08_04       1436 non-null   int64  
 4   KM              1436 non-null   int64  
 5   HP              1436 non-null   int64  
 6   Doors           1436 non-null   int64  
 7   Gears           1436 non-null   int64  
 8   Weight          1436 non-null   int64  
dtypes: int64(8), object(1)
memory usage: 101.1+ KB
```

```
In [32]: cars2=pd.concat([cars1.iloc[:,0:]],axis=1)
cars2
```

Out[32]:

	Price	Age_08_04	KM	HP	Doors	Cylinders	Gears	Weight
0	13500	23	46986	90	3	4	5	1165
1	13750	23	72937	90	3	4	5	1165
2	13950	24	41711	90	3	4	5	1165
3	14950	26	48000	90	3	4	5	1165
4	13750	30	38500	90	3	4	5	1170
...	...	...	...	...	...	...	...	...
1431	7500	69	20544	86	3	4	5	1025
1432	10845	72	19000	86	3	4	5	1015
1433	8500	71	17016	86	3	4	5	1015
1434	7250	70	16916	86	3	4	5	1015
1435	6950	76	1	110	5	4	5	1114

1436 rows × 8 columns

```
In [39]: cars3=cars2.rename({'Age_08_04':'Age'},axis=1)
cars3
```

Out[39]:

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
0	13500	23	46986	90	3	4	5	1165
1	13750	23	72937	90	3	4	5	1165
2	13950	24	41711	90	3	4	5	1165
3	14950	26	48000	90	3	4	5	1165
4	13750	30	38500	90	3	4	5	1170
...	...	...	...	...	...	...	...	...
1431	7500	69	20544	86	3	4	5	1025
1432	10845	72	19000	86	3	4	5	1015
1433	8500	71	17016	86	3	4	5	1015
1434	7250	70	16916	86	3	4	5	1015
1435	6950	76	1	110	5	4	5	1114

1436 rows × 8 columns

```
In [40]: cars3[cars3.duplicated()]
```

Out[40]:

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
113	24950	8	13253	116	5	4	5	1320

```
In [41]: cars3=cars3.drop_duplicates().reset_index(drop=True)
cars3
```

Out[41]:

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
0	13500	23	46986	90	3	4	5	1165
1	13750	23	72937	90	3	4	5	1165
2	13950	24	41711	90	3	4	5	1165
3	14950	26	48000	90	3	4	5	1165
4	13750	30	38500	90	3	4	5	1170
...	...	...	...	...	...	...	...	...
1430	7500	69	20544	86	3	4	5	1025
1431	10845	72	19000	86	3	4	5	1015
1432	8500	71	17016	86	3	4	5	1015
1433	7250	70	16916	86	3	4	5	1015
1434	6950	76	1	110	5	4	5	1114

1435 rows × 8 columns

```
In [42]: cars3
```

Out[42]:

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
0	13500	23	46986	90	3	4	5	1165
1	13750	23	72937	90	3	4	5	1165
2	13950	24	41711	90	3	4	5	1165
3	14950	26	48000	90	3	4	5	1165
4	13750	30	38500	90	3	4	5	1170
...	...	...	...	...	...	...	...	...
1430	7500	69	20544	86	3	4	5	1025
1431	10845	72	19000	86	3	4	5	1015
1432	8500	71	17016	86	3	4	5	1015
1433	7250	70	16916	86	3	4	5	1015
1434	6950	76	1	110	5	4	5	1114

1435 rows × 8 columns

In [43]: cars3.describe()

Out[43]:

	Price	Age	KM	HP	Doors	Cylinders	Gears
count	1435.000000	1435.000000	1435.000000	1435.000000	1435.000000	1435.0	1435.000000
mean	10720.915679	55.980488	68571.782578	101.491986	4.032753	4.0	5.026481
std	3608.732978	18.563312	37491.094553	14.981408	0.952667	0.0	0.188575
min	4350.000000	1.000000	1.000000	69.000000	2.000000	4.0	3.000000
25%	8450.000000	44.000000	43000.000000	90.000000	3.000000	4.0	5.000000
50%	9900.000000	61.000000	63451.000000	110.000000	4.000000	4.0	5.000000
75%	11950.000000	70.000000	87041.500000	110.000000	5.000000	4.0	5.000000
max	32500.000000	80.000000	243000.000000	192.000000	5.000000	4.0	6.000000



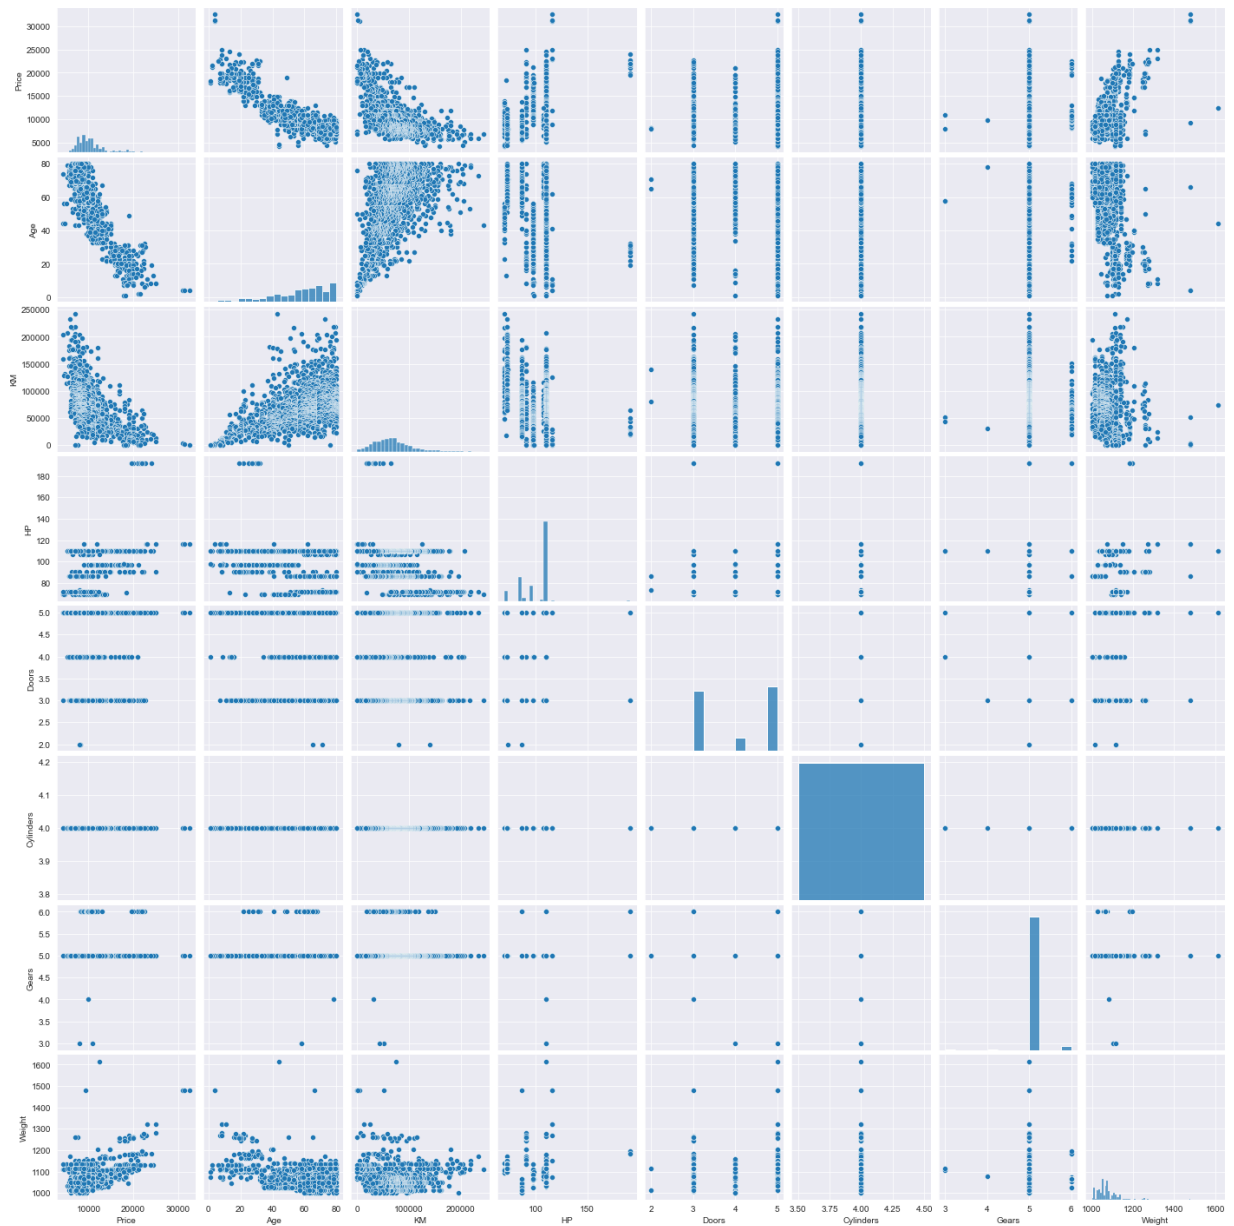
In [44]: cars3.corr()

Out[44]:

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
Price	1.000000	-0.876273	-0.569420	0.314134	0.183604	NaN	0.063831	0.575869
Age	-0.876273	1.000000	0.504575	-0.155293	-0.146929	NaN	-0.005629	-0.466484
KM	-0.569420	0.504575	1.000000	-0.332904	-0.035193	NaN	0.014890	-0.023969
HP	0.314134	-0.155293	-0.332904	1.000000	0.091803	NaN	0.209642	0.087143
Doors	0.183604	-0.146929	-0.035193	0.091803	1.000000	NaN	-0.160101	0.301734
Cylinders	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Gears	0.063831	-0.005629	0.014890	0.209642	-0.160101	NaN	1.000000	0.021238
Weight	0.575869	-0.466484	-0.023969	0.087143	0.301734	NaN	0.021238	1.000000

```
In [45]: sns.set_style(style='darkgrid')
sns.pairplot(cars3)
```

```
Out[45]: <seaborn.axisgrid.PairGrid at 0x2c7b2c7c970>
```



```
In [46]: model=snf.ols('Price~Age+KM+HP+Doors+Gears+Weight',data=cars3).fit()
```

```
In [47]: model.params
```

```
Out[47]: Intercept    -6838.987234  
Age                -122.288250  
KM                  -0.019928  
HP                  28.327782  
Doors               -8.715826  
Gears               625.297840  
Weight              18.455133  
dtype: float64
```

```
In [48]: model.tvalues , np.round(model.pvalues,5)
```

```
Out[48]: (Intercept    -5.204801  
Age                -46.774676  
KM                 -16.490150  
HP                 10.840831  
Doors              -0.218061  
Gears               3.169455  
Weight             22.141591  
dtype: float64,  
Intercept     0.00000  
Age           0.00000  
KM            0.00000  
HP            0.00000  
Doors         0.82741  
Gears         0.00156  
Weight        0.00000  
dtype: float64)
```

```
In [49]: model.rsquared , model.rsquared_adj
```

```
Out[49]: (0.8615946984866649, 0.8610131636063568)
```

```
In [50]: slr_c=snf.ols('Price~Doors',data=cars3).fit()  
slr_c.tvalues , slr_c.pvalues
```

```
Out[50]: (Intercept    19.421546  
Doors         7.070520  
dtype: float64,  
Intercept     8.976407e-75  
Doors         2.404166e-12  
dtype: float64)
```

In [51]:

```
rsq_age=snf.ols('Age~KM+HP+Doors+Gears+Weight',data=cars3).fit().rsquared
vif_age=1/(1-rsq_age)

rsq_KM=snf.ols('KM~Age+HP+Doors+Gears+Weight',data=cars3).fit().rsquared
vif_KM=1/(1-rsq_KM)

rsq_HP=snf.ols('HP~Age+KM+Doors+Gears+Weight',data=cars3).fit().rsquared
vif_HP=1/(1-rsq_HP)

rsq_DR=snf.ols('Doors~Age+KM+HP+Gears+Weight',data=cars3).fit().rsquared
vif_DR=1/(1-rsq_DR)

rsq_GR=snf.ols('Gears~Age+KM+HP+Doors+Weight',data=cars3).fit().rsquared
vif_GR=1/(1-rsq_GR)

rsq_WT=snf.ols('Weight~Age+KM+HP+Doors+Gears',data=cars3).fit().rsquared
vif_WT=1/(1-rsq_WT)


# Putting the values in Dataframe format
d1={'Variables':['Age','KM','HP','Doors','Gears','Weight'],
    'Vif':[vif_age,vif_KM,vif_HP,vif_DR,vif_GR,vif_WT]}
Vif_df=pd.DataFrame(d1)
Vif_df
```

Out[51]:

	Variables	Vif
0	Age	1.866057
1	KM	1.626264
2	HP	1.214147
3	Doors	1.148708
4	Gears	1.096575
5	Weight	1.502749



```
In [58]: sm.qqplot(model.resid,line='q') # 'q' - A line is fit through the quantiles # Lin
plt.title("Normal Q-Q plot of residuals")
plt.show()
```



```
In [71]: list(np.where(model.resid<-6000))
```


```
Out[71]: [array([220, 600, 959], dtype=int64)]
```

```
In [72]: list(np.where(model.resid>6000))
```

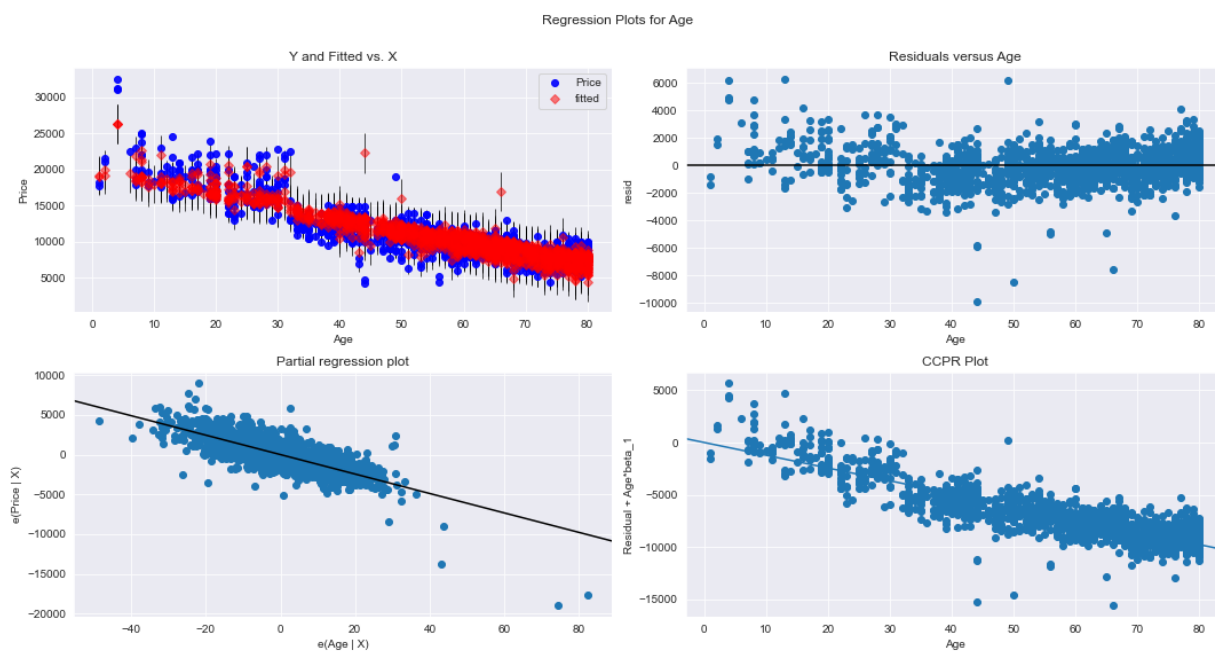
```
Out[72]: [array([109, 146, 522], dtype=int64)]
```

```
In [73]: def standard_values(vals) : return (vals-vals.mean())/vals.std()
```

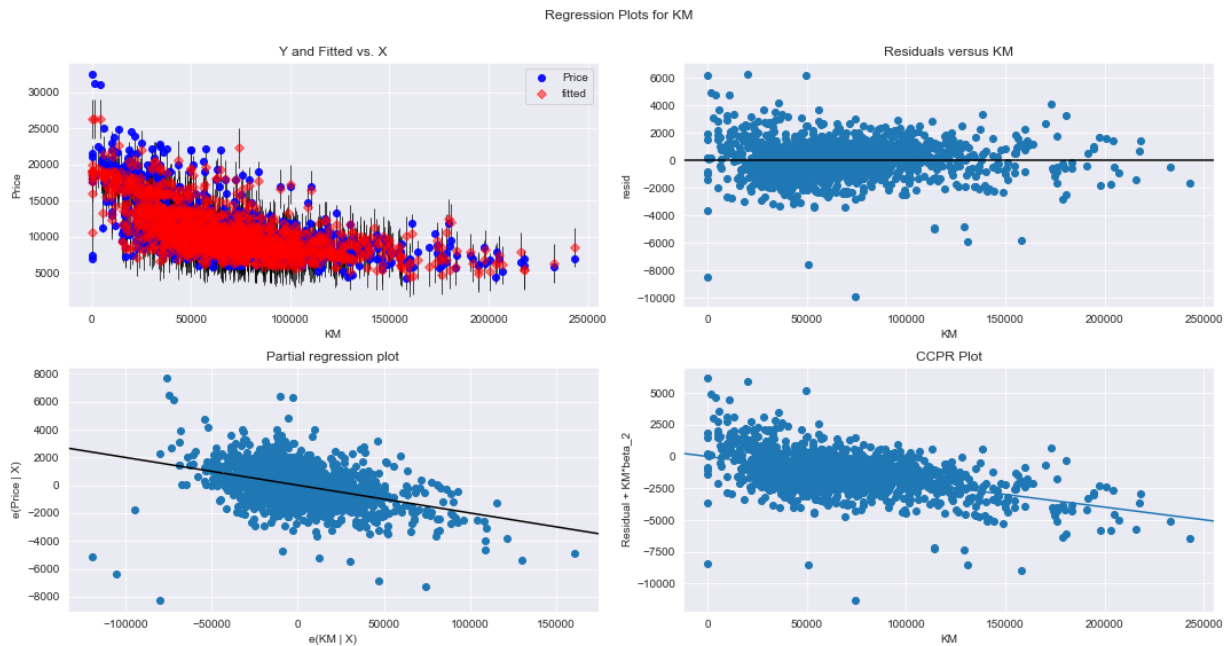
```
In [74]: plt.scatter(standard_values(model.fittedvalues),standard_values(model.resid))
plt.title('Residual Plot')
plt.xlabel('standardized fitted values')
plt.ylabel('standardized residual values')
plt.show()
```



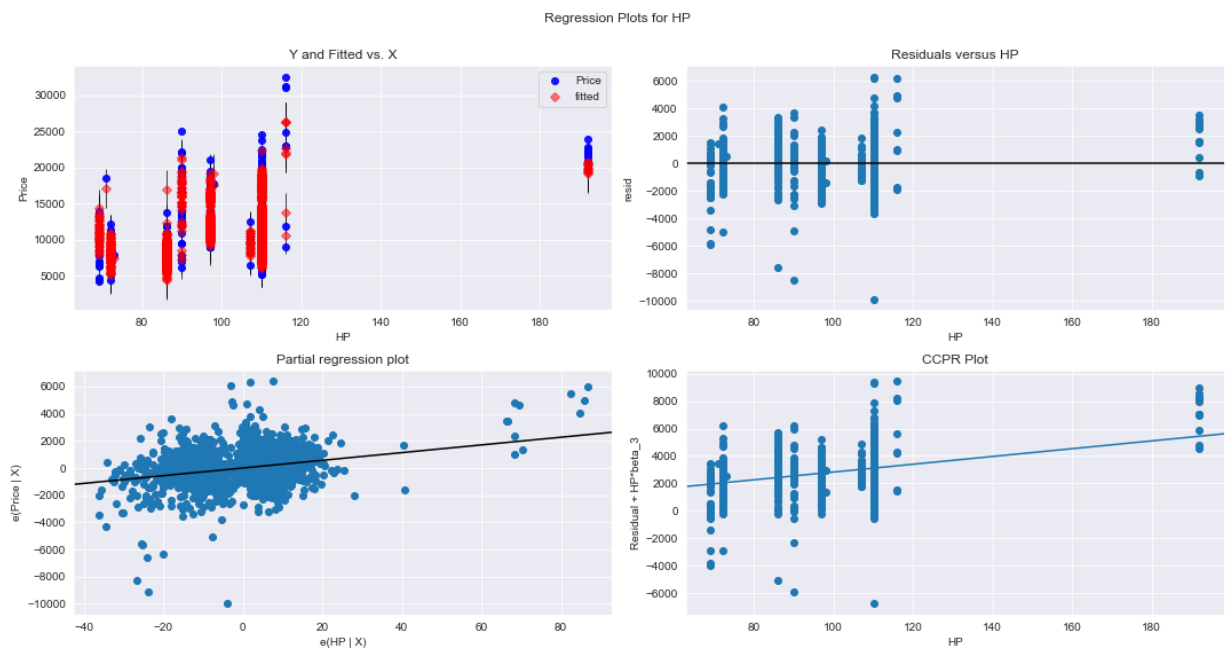
```
In [76]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Age',fig=fig)
plt.show()
```



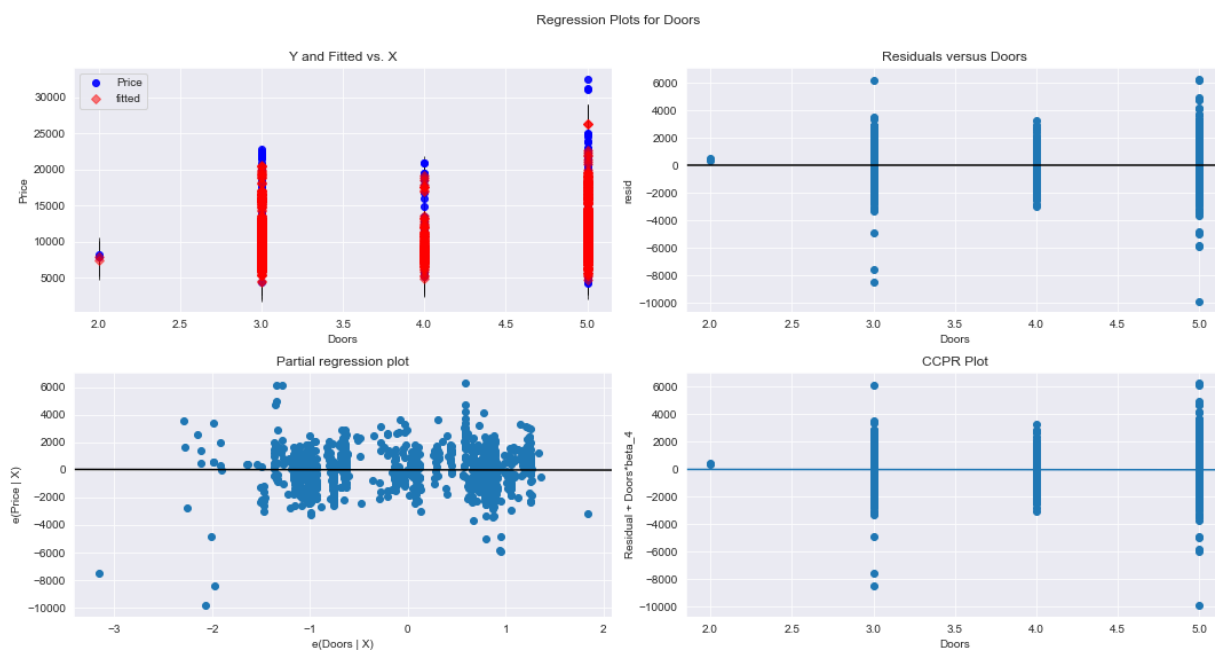
```
In [77]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'KM',fig=fig)
plt.show()
```



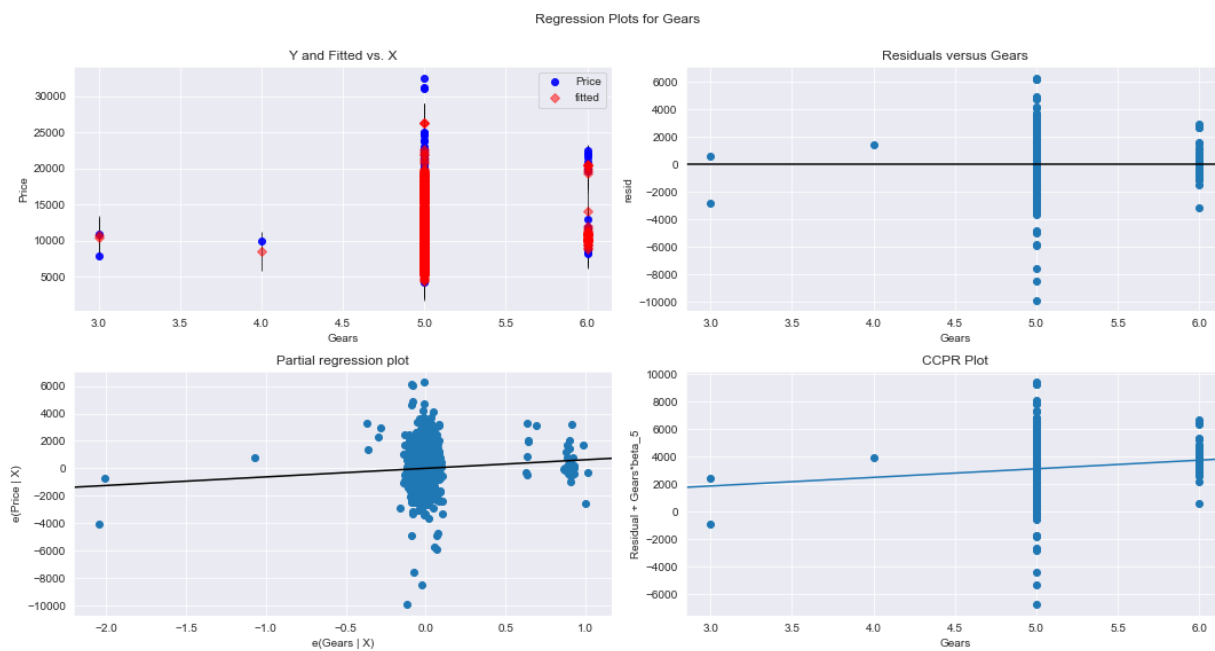
```
In [56]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'HP',fig=fig)
plt.show()
```



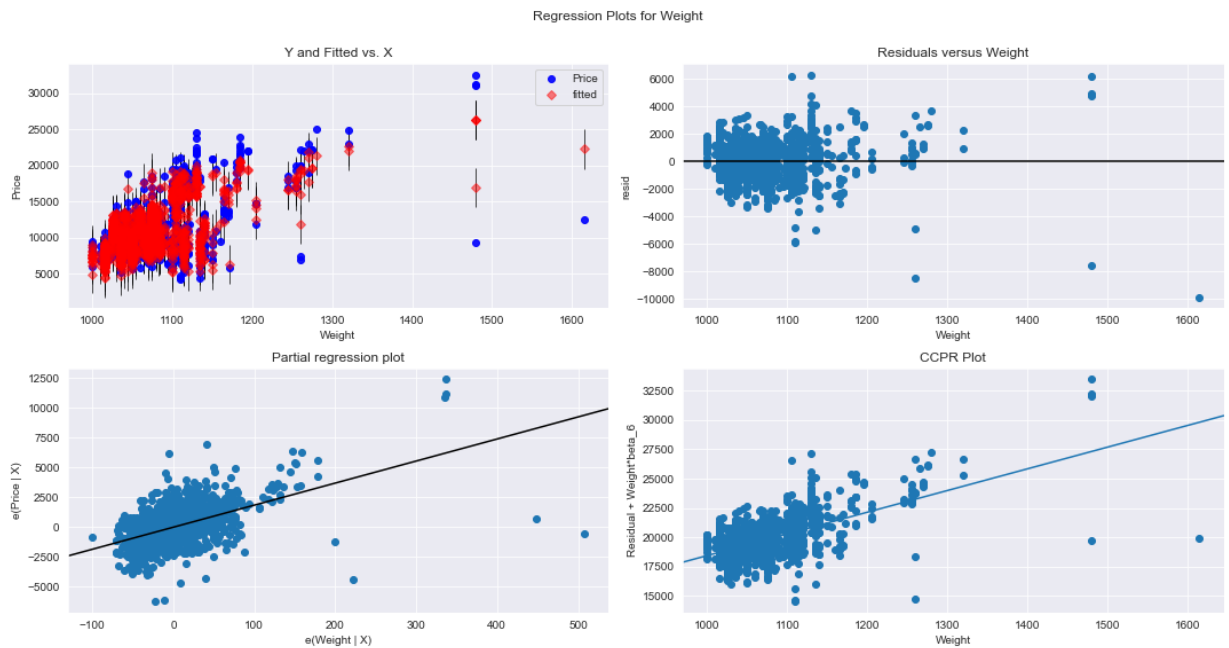
```
In [78]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'Doors',fig=fig)
plt.show()
```



```
In [64]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'Gears',fig=fig)
plt.show()
```




```
In [79]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'Weight',fig=fig)
plt.show()
```



```
In [80]: (c,_)=model.get_influence().cooks_distance
c
```

```
Out[80]: array([4.84834865e-03, 2.81504747e-03, 3.49062601e-03, ...,
                4.29681961e-06, 8.15626746e-04, 1.20038984e-02])
```

```
In [66]: fig=plt.figure(figsize=(20,7))
plt.stem(np.arange(len(cars3)),np.round(c,3))
plt.xlabel('Row Index')
plt.ylabel('Cooks Distance')
plt.show()
```



```
In [81]: np.argmax(c) , np.max(c)
```

Out[81]: (220, 0.9561392473392505)

```
In [82]: fig,ax=plt.subplots(figsize=(20,20))
fig=influence_plot(model,ax = ax)
```



```
In [84]: k=cars3.shape[1]
n=cars3.shape[0]
leverage_cutoff = (3*(k+1))/n
leverage_cutoff
```

Out[84]: 0.018815331010452963

```
In [93]: cars3[cars3.index.isin([220])]
```

Out[93]:

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
220	12450	44	74172	110	5	4	5	1615

```
In [86]: cars4=cars3.copy()  
cars4
```

Out[86]:

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
0	13500	23	46986	90	3	4	5	1165
1	13750	23	72937	90	3	4	5	1165
2	13950	24	41711	90	3	4	5	1165
3	14950	26	48000	90	3	4	5	1165
4	13750	30	38500	90	3	4	5	1170
...	...	...	...	...	...	...	...	...
1430	7500	69	20544	86	3	4	5	1025
1431	10845	72	19000	86	3	4	5	1015
1432	8500	71	17016	86	3	4	5	1015
1433	7250	70	16916	86	3	4	5	1015
1434	6950	76	1	110	5	4	5	1114

1435 rows × 8 columns

```
In [94]: cars5=cars4.drop(cars4.index[[220]],axis=0).reset_index(drop=True)
cars5
```

Out[94]:

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
0	13500	23	46986	90	3	4	5	1165
1	13750	23	72937	90	3	4	5	1165
2	13950	24	41711	90	3	4	5	1165
3	14950	26	48000	90	3	4	5	1165
4	13750	30	38500	90	3	4	5	1170
...	...	...	...	...	...	...	...	...
1429	7500	69	20544	86	3	4	5	1025
1430	10845	72	19000	86	3	4	5	1015
1431	8500	71	17016	86	3	4	5	1015
1432	7250	70	16916	86	3	4	5	1015
1433	6950	76	1	110	5	4	5	1114

1434 rows × 8 columns

```
In [95]: while np.max(c)>0.5 :
          model=snf.ols('Price~Age+KM+HP+Doors+Gears+Weight',data=cars5).fit()
          (c,_)=model.get_influence().cooks_distance
          c
          np.argmax(c) , np.max(c)
          cars5=cars5.drop(cars5.index[[np.argmax(c)]],axis=0).reset_index(drop=True)
          cars5
        else:
          final_model=snf.ols('Price~Age+KM+HP+Doors+Gears+Weight',data=cars5).fit()
          final_model.rsquared , final_model.aic
          print("Thus model accuracy is improved to",final_model.rsquared)
```

Thus model accuracy is improved to 0.8673586781804876

```
In [96]: final_model.rsquared
```

Out[96]: 0.8673586781804876



```
In [97]: cars5
```

```
Out[97]:
```

	Price	Age	KM	HP	Doors	Cylinders	Gears	Weight
0	13500	23	46986	90	3	4	5	1165
1	13750	23	72937	90	3	4	5	1165
2	13950	24	41711	90	3	4	5	1165
3	14950	26	48000	90	3	4	5	1165
4	13750	30	38500	90	3	4	5	1170
...	...	...	...	...	...	...	...	...
1429	7500	69	20544	86	3	4	5	1025
1430	10845	72	19000	86	3	4	5	1015
1431	8500	71	17016	86	3	4	5	1015
1432	7250	70	16916	86	3	4	5	1015
1433	6950	76	1	110	5	4	5	1114

1434 rows × 8 columns

```
In [98]: new_data=pd.DataFrame({'Age':12,"KM":40000,"HP":80,"Doors":4,"Gears":5,"QT":69,"v":1012})
new_data
```

```
Out[98]:
```

	Age	KM	HP	Doors	Gears	QT	Weight
0	12	40000	80	4	5	69	1012

```
In [99]: final_model.predict(new_data)
```

```
Out[99]: 0    14699.298039
dtype: float64
```

```
In [100]: pred_y=final_model.predict(cars5)
pred_y
```

```
Out[100]: 0      16704.446380
1      16168.971437
2      16694.086877
3      16325.910639
4      16148.065524

...
1429    8771.451048
1430    8239.800877
1431    8399.943149
1432    8521.210819
1433   10811.940431
Length: 1434, dtype: float64
```